

Project Instructions

I. Required Software Setup

Nachos Code Distributions

The Java version of Nachos is available for you to download from learn.tsinghua.edu.cn (Under course files filenames are nachos-java.tar.gz). You can unpack this file using the command

```
gzcat nachos-java.tar.gz | tar xf -
```

which will create the directory nachos. See the file README contained within the tarfile for information on using and compiling the code.

Each Subversion repository comes with the recommended repository structure and the Nachos distribution checked in. You should be reading through the code and trying to understand how it works.

Java Distributions

The Java version of Nachos requires Java SE Java Development Kit 1.5 or later. You can download the appropriate Java SE JDK for your machine from here: <http://www.oracle.com/technetwork/java/javase/downloads/index.html>.

MIPS Cross-compilers

Nachos executes user applications on a simulated MIPS computer. In order to compile these programs, you need a MIPS cross-compiler. You must download and install the appropriate cross-compiler from <https://inst.eecs.berkeley.edu/~cs162/sp05/Nachos/xgcc.shtml>.

II. General Project Information

There are four projects. Each project has three components:

1. A design document.
2. Project code.
3. Project group member evaluations.

Project Grading

The design document will be worth 20% of the project grade, and the code itself will be worth the

other 80%.

Project Deadlines

You should submit assignments as early as possible to avoid submission problems.

Assignments (design docs, code, and group evaluations) are due by 11:59pm on the due date.

Submission Instructions

Submission instructions for the design doc and source code in project 1

- Only one member in the group needs to submit the homework, to minimize prospective confusions, please only designate one member to submit the homework for each group assignment. We do deduct points for multiple submissions – for poor team organization and communication.
- Please submit the following files to learn.tsinghua.edu.cn
 - Put your group number and names in the text box
 - Your design document, please name it as proj1-design.pdf
 - Your source code as a tar file (or zip file), generated by the “submit” script, as separately described in each of the four projects.
- You can submit as many times as you want to before the deadline, only the last submission before the deadline counts.

Project Deadlines

- Project #1: Threads
 - Code due 4/4
- Project #2: Multiprogramming
 - Code due 4/25
- Project #3: Single node key-value store (Client-Server Communication)
 - Code due 5/23
- Project #4: TBD
 - Code due TBD

III. Design document instructions

Overview

- Be sure to read through this entire document. It's all relevant.
- Each design document is worth 20% of the project. While it will likely take less than 20% of the time you spend on the project, you should take it very seriously.
- We will grade your designs harshly. The design is essentially the most important part of the project. Having a good project design can literally cut your total coding time by a factor of 10.
- Design documents should be around 2,000 to 4,000 English words long. If it's longer than 5,000 words, we won't read it. So keep'em short and to the point.

- Design documents will be submitted online, in PDF or text format. We encourage you to use your favorite word processor to make your design document, but you must convert it to PDF or text when you're done. We'll provide basic directions on how to do that.

What goes into a design document?

A design document is a complete high-level solution to the problem presented. It should be detailed enough that somebody who already understands the problem could go out and code the project without having to make any significant decisions. Further, if this somebody happens to be an experienced coder, they should be able to use the design document to code the solution in a few hours (not necessarily including debugging).

So what actually goes in?

A high-level description of your solution, including **design decisions** and **data structures**

Declarations for all new classes, procedures, and global/class variables

Descriptions of all new procedures (unless you can tell exactly what it does from the name), including the purpose of the procedure, and an explanation of how it works and/or pseudocode

For example, this is the pseudocode I would write for the existing `Condition::Wait`:

```
void Condition::Wait()
{
    waiter = new Semaphore, initially 0
    append waiter to waitQueue
    conditionLock->Release()
    waiter->P()
    conditionLock->Acquire()
}
```

Your pseudocode has to be precise. For instance, in describing your solution for the Communicator class, it is not enough to say

```
if anybody's waiting, then signal cv
```

You need to say how you determine if anybody's waiting, e.g. by saying

```
if (waitingReceivers>0 or waitingSenders>0) then signal cv
```

Also, another important thing to remember is that a design document needs to include the **correctness invariants** and **testing strategy**. The testing strategy includes a clear plan of the testing methodology and may include a description of test cases that will be used to test correctness invariants. Focus on the testing strategy. If you want, you may itemize things that will need testing.

What *doesn't* go into a design document?

Keep in mind that your TA understands the project very well. Do not restate the problem in your design document. Your TA is far more interested in your solution than in knowing that you understood the problem.

IV. Guide to Nachos

Nachos Project instructions (first two projects)

The first two projects for this course is to build an operating system from scratch. We believe that the best way for you to understand operating systems concepts is to build a real operating system and then to experiment with it. To help you get started, people have built a very simple, but functional, operating system called Nachos. Over the course of the semester, your job will be to improve the functionality and performance of Nachos; we expect that you will eventually modify or replace much of what we have written.

The project has two phases, corresponding to two of the major pieces of a modern operating system: thread management, and multiprogramming.

Some phases build on previous phases; for example, all parts of the project use thread management routines. As far as possible, we have structured the assignments so that you will be able to finish the project even if all of the pieces are not working, but you will get more out of the project if you use your own software.

Part of the code we provide is a software emulation of a network of MIPS-like workstations. For instance, our code emulates turning on and off interrupts, taking exceptions and page faults, as well as the actions of physical devices (e.g. a disk, a console, and a network controller).

It is important to realize that while we run Nachos on top of this emulation as a Java program, the code you write and most of the code we provide are exactly the same as if Nachos were running on bare hardware. We run Nachos as a Java program for convenience: it is very portable, it is easier to debug (Java is a type-safe language, and we can use jdb), and so that multiple students can run Nachos at the same time on the same machine. These same reasons apply in industry; it is usually a good idea to test out system code in a simulated environment before running it on potentially flaky hardware.

In order to port Nachos to real hardware, we would have to replace our emulation code with a little bit of machine-dependent code and some physical machines. For example, in phase 1, we provide routines to enable and disable interrupts; on real hardware, the CPU provides these functions.

Unless you work for a really smart company, when you develop operating system software you usually cannot change the hardware to make your life easier. Thus, you are not permitted to change any of our emulation code, although you are permitted to change most of the Nachos code that runs on top of the emulation.

Finally, some suggestions for doing well in this course:

Read the code that is handed out with the assignment first, until you pretty much understand it. Start with the interface documentation.

Don't code until you understand what you are doing. Design, design, and design first. Only then split up what each of you will code.

If you get stuck, talk to as many other people as possible.

[A guide for Nachos](#)

Will be handed out in class. Also posted on learn.tsinghua.edu.cn